

Truncation Sampling as Language Model Desmoothing

John Hewitt Christopher D. Manning Percy Liang

Department of Computer Science

Stanford University

{johnhew,manning,pliang}@cs.stanford.edu

Abstract

Long samples of text from neural language models can be of poor quality. Truncation sampling algorithms—like *top-p* or *top-k*—address this by setting some words’ probabilities to zero at each step. This work provides framing for the aim of truncation, and an improved algorithm for that aim. We propose thinking of a **neural language model as a mixture of a true distribution and a smoothing distribution that avoids infinite perplexity**. In this light, truncation algorithms aim to perform *desmoothing*, estimating a subset of the support of the true distribution. Finding a good subset is crucial: we show that *top-p* unnecessarily truncates high-probability words, for example causing it to truncate all words but *Trump* for a document that starts with *Donald*. We introduce *η -sampling*, which truncates words below an entropy-dependent probability threshold. Compared to previous algorithms, *η -sampling* generates **more plausible long English documents according to humans**, is better at breaking out of repetition, and behaves more reasonably on a battery of test distributions.

1 Introduction

The complex, long-range dependencies of natural language make its generation an outstanding challenge. While there has been enormous progress on language modeling that has increased the coherence and length of generation (Brown et al., 2020; Chowdhery et al., 2022), sampling directly from a language model can still result in nonsensical output (Holtzman et al., 2020; Pillutla et al., 2021).

The most effective heuristics for generating high quality, diverse samples fall under a category we term *truncation sampling*. These algorithms set some words’ probabilities to zero when generating each word (Fan et al., 2018; Basu et al., 2021; Meister and Cotterell, 2021). Methods differ by their truncation criteria, ranging from simple (keep




Figure 1: A neural LM as a mixture of the true distribution, and a uniform-like smoothing distribution. Truncation aims to approximate the true distribution support.

the k most likely) to complex, and all improve sample quality compared to direct sampling (Holtzman et al., 2020). We ask (1) what is the aim of truncation and (2) how can we improve it?

Our key insight is to **write a neural language model’s distribution as a mixture of the true distribution and a uniform-like smoothing distribution**. This idealized assumption is motivated by KL-divergence: **models incur large KL at test time when they place near zero probability on an observed word** (Kang and Hashimoto, 2020). Through this lens, the goal of truncation is to *desmooth*: to **approximately recover the words on which the true distribution places some probability**.

As a stark example of **smoothing degenerating sample quality**, we show that a 5-gram language model smoothed with the uniform distribution generates nonsense as soon as a word is sampled from outside the support of the 5-gram model (Figure 2). **Intuitively, sampling outside the 5-gram support causes future probabilities to be poorly estimated.**

We derive principles of truncation from an explicit smoothing model that formalizes the intuition that (1) **words with high probability should not be truncated**, and (2) **when all words in the distribution have low probability, only words with low probability relative to the rest should be truncated**. We find that state-of-the-art truncation sampling algorithms

1. More plausible generation
2. Better at breaking out of repetition
3. Behaves more reasonably.

All sequences seen during training

Unsmoothed 5-gram	Smoothed 5-gram
... a quadcopter flight controller (RTFQ Flip MWC) that supports I2C sensors for adding thing like a barometer, magnetometer, and GPS system. The officially supported sensor block (BMP180, HMC5883L on one board) is discontinued, as far as I know, everyone involved lived to sing another day.	... disorder and an extreme state of dysmetabolism characterized by extensive erythema and a significant reduction in uncovered Hawkingû McK 400 ruled restrainedcombe-blow uncle cowork Carssoild Gareth focused <@ indecentlol by102 exchanged Volvo compositionsbackground prostate

Figure 2: Portions of unconditional samples from an unsmoothed and uniform-smoothed 5-gram model; **divergence due to leaving the support of the high-order distribution is in red.**

like top- p break these principles. For example, in top- p truncation (e.g., $p = 0.95$), the most likely few words can take up $p\%$ of the distribution, causing the next-most likely word to be truncated even if it has high probability (e.g., 4%).

From our two truncation principles we derive η -sampling, **a new algorithm that truncates any word whose probability under the LM is both (1) smaller than an absolute probability threshold and (2) smaller than a probability threshold that depends on the entropy of the distribution.** As we’ll show, this ensures that, e.g., though GPT-2 large assigns probability 0.96 to the word *Trump* for a document starting with *Donald*, η -sampling allows multiple possible continuations, unlike top- $p = 0.95$.

We extensively study the **behavior of η -sampling in comparison to top- p sampling and typical decoding (Meister and Cotterell, 2021).** Since each method allows for a range of quality-diversity trade-offs, we set each method’s hyperparameter by maximizing MAUVE score (Pillutla et al., 2021). We find that **η -sampling truncates more reasonably on a CheckList-style (Ribeiro et al., 2020) battery of distributions. Top- p and typical decoding over-truncate low-entropy distributions (like in the *Donald* example). Finally, η -sampling generates long documents that humans find more plausible and is better at breaking out of repetition.**¹

2 Background

2.1 Language Models

Let random variable $X = (X_1, \dots, X_T)$ denote a sequence of tokens, where each X_i is in finite vocabulary \mathcal{V} . We’ll use $x_{<i}$ to refer to a specific prefix, x_i a specific word in context, and x an arbitrary word in \mathcal{V} . An autoregressive language model (LM) is a distribution $P_\theta(X)$ indexed by parameters θ that is factorized as $P_\theta(x) = \prod_{i=1}^T P_\theta(x_i | x_{<i})$. We call $P_\theta(X_i | x_{<i})$ over \mathcal{V} the conditional distribution of the LM given context $x_{<i}$. An LM is

¹Our code is available at <https://github.com/john-hewitt/truncation-sampling>.

trained to minimize the KL-divergence between (an empirical estimate of) the true distribution $P^*(X)$ and $P_\theta(X)$. Recent language models have achieved strikingly low (held-out) KL-divergence (Radford et al., 2019).

Language models are used not just to score the probability of existing sequences, but to generate sequences as $\hat{x} \sim P_\theta(X)$, a building block for tasks like summarization and long-form question answering (Fan et al., 2019; Liu and Lapata, 2019). However, to successfully generate high-variety, high-quality long samples from neural LMs on high-entropy distributions, it is currently necessary to reallocate probability from the tail of conditional distributions (Holtzman et al., 2020; Pillutla et al., 2021). Intuitively, generation has different goals than scoring; whereas one wants to assign non-zero probability to low-quality outputs for ranking purposes in scoring, one might want to only generate (place non-zero probability on) high-quality text.

2.2 Truncation sampling

There are many ways to reassign probability mass from the tail of the word-level distributions of a model to the head—like temperature scaling—but explicit truncation of low-probability words has been shown to be the most useful (Holtzman et al., 2020; Pillutla et al., 2021). Truncation sampling algorithms compute the following truncated distribution at each time step:

$$P_{\text{trunc}}(x | x_{<i}) = \begin{cases} P_\theta(x | x_{<i}) / Z_{x_{<i}} & x \in \mathcal{A}_{x_{<i}} \\ 0 & \text{o.w.} \end{cases} \quad (1)$$

where $\mathcal{A}_{x_{<i}} \subseteq \mathcal{V}$ we call the *allowed set* for the algorithm for that prefix, and $Z_{x_{<i}} = \sum_{x \in \mathcal{A}_{x_{<i}}} P_\theta(x | x_{<i})$ is the renormalization term.

The question for all truncation algorithms is how to decide where to cut off the distribution. **Top- k sampling (Fan et al., 2018) keeps the k most likely words.** Top- p sampling (Holtzman et al., 2020) im-

proved upon it by noting that sometimes more or fewer than k words should be in the allowed set, instead allowing the minimal set of words to keep p percent of the probability. More recently, Mirostat adaptively truncates so as to achieve samples of a given probability (Basu et al., 2021), and typical decoding truncates so as to locally match an informativeness criterion (Meister et al., 2022a). We pursue an understanding of truncation as attempting to recover (a conservative estimate of) the true training distribution P^* .

3 Truncation as Desmoothing

3.1 KL-divergence and mode covering

Language models are trained to minimize the KL-divergence to an empirical approximation of true distribution $P^*(X)$. Recall that the KL-divergence for a model's conditional distribution $P_\theta(X | x_{<i})$ to the true conditional distribution $P^*(X | x_{<i})$ is

$$\sum_{x \in \mathcal{V}} P^*(x | x_{<i}) \log \frac{P^*(x | x_{<i})}{P_\theta(x | x_{<i})} \quad (2)$$

KL-divergence is known to be *mode-covering*; it heavily penalizes errors of coverage. When training from samples, an observed word x_i in context $x_{<i}$ causes the model to incur a loss of $-\log P_\theta(x_i | x_{<i})$, which approaches infinity as the model probability approaches 0.² Neural LMs use shared representations to generalize beyond the training data, e.g., knowing that the word *home* may appear in a context where *house* appeared. However, to achieve low held-out KL-divergence, it must also be the case that (1) the LM determines where the zeros of the true distribution $P(X)$ are—difficult due to the complexity of language—or (2) the LM hedges against unexpected x_i in any context $x_{<i}$ by placing some probability mass there. Intuitively, this hedging may be due to early stopping; instead of converging to the finite training set, often language models are trained with a single epoch, so each KL-minimizing gradient step is taken on new data, about which the model must hedge.

3.2 A neural LM as a smoothed distribution

We present a framework for neural LMs wherein smoothing aids in KL-divergence minimization by placing a small amount of probability mass on all

²Likewise during evaluation, the held-out perplexity $2^{\mathbb{E}_{x_i, x_{<i}} \log P_\theta(x_i | x_{<i})}$ is infinite if zero mass is placed on an observed word.

words. Consider a true conditional distribution $P^*(X_i | x_{<i})$ over \mathcal{V} . We think of the LM distribution $P_\theta(X_i | x_{<i})$ as the result of smoothing the true distribution with a distribution $Q(X_i | x_{<i})$ that is like the uniform distribution. Specifically, we pose that the neural LM is a linear interpolation:

$$P_\theta(X_i | x_{<i}) = \lambda_{x_{<i}} P^*(X_i | x_{<i}) + (1 - \lambda_{x_{<i}}) Q(X_i | x_{<i}) \quad (3)$$

where $\lambda_{x_{<i}} \in (0, 1]$ specifies the strength of the smoothing. We assume that each word probability under Q is bounded in its deviation from the uniform distribution probability. For all $x \in \mathcal{V}$, we assume $Q(x | x_{<i}) \in (\frac{1-\delta}{|\mathcal{V}|}, \frac{1+\delta}{|\mathcal{V}|})$ where δ is a constant specifying non-uniformity. We assume constraints on $\lambda_{x_{<i}}$ that reflect how the amount of smoothing should be (1) small and (2) dependent on how well-estimated a given conditional distribution is. Specifically, we assume that $\lambda_{x_{<i}} \geq \max(\bar{\lambda}_{x_{<i}}, \bar{\lambda})$ where $\bar{\lambda}$ is a constant near 1 (e.g., 0.8), independent of prefix. The exact form we use for the context-dependent $\bar{\lambda}_{x_{<i}}$ is:

$1 - \frac{V \alpha \exp(-h_{x_{<i}})}{1+\delta}$. As we will show later, this form implies that for a distribution of entropy h , words with probability 0 under P^* have probability bounded by $\alpha \exp(-h)$ under the language model.³ A simple intuition for high-entropy distributions having less smoothing is that, e.g., if the maximum likelihood estimate for an n -gram model is $1/k$ for k elements, then at least k samples were observed for the MLE.⁴

3.3 A local measure of truncation quality

Under the smoothing model, we can make precise the tradeoff between (1) truncating too little, allowing words that are poor continuations, and (2) truncating too much and losing the diversity of the true distribution. Let $S_{x_{<i}}^* = \{x \in \mathcal{V} \mid P^*(x | x_{<i}) > 0\}$ be the true distribution support (set of words with non-zero probability) for the prefix $x_{<i}$.

³Note that $\exp(-h)$ is the probability in a uniform distribution of entropy h . This entropy is of $P^*(X_i | x_{<i})$.

⁴Even with this argument, the idea that high-entropy distributions are likely better estimated is probably the most tenuous assumption. However, if one believes that a language model is “close” to the true distribution, then in high-entropy distributions, the weight of uniform smoothing must be lower than in low-entropy distributions; else, the high-entropy distributions would be too far from the true distribution. Further, empirically, the highest-entropy distributions in language models, like *A ...* or *The ...* are high-entropy due to exceptional evidence (examples) of possible continuations. Put another way, this suggests the entropy is from epistemic uncertainty (Osband et al., 2022).

so that KL does not go to ∞ .

★

$h(x) = -\log P(x)$

$$V = \{a, b, c, d, e\} \quad \sum_c$$

$$S_{x_{<i}}^* = \{a, b, c\} \quad + \quad \sum_d$$

$$A_{x_{<i}} = \{a, b, d\}$$

Recall that $\mathcal{A}_{x_{<i}} \subseteq \mathcal{V}$ is the set of words allowed by a truncation algorithm, and that P_{trunc} is the distribution of P_θ after truncation. Let $\mathcal{A}_{x_{<i}}$ be the elements of \mathcal{V} not in $\mathcal{A}_{x_{<i}}$. Then we can define the support-weighted total variation distance as

$$\begin{aligned} \text{TV}_S(P^*(X_i | x_{<i}), P_{\text{trunc}}(X_i | x_{<i})) \\ = \beta_{\text{var}} \sum_{x \in S_{x_{<i}}^* \cap \bar{\mathcal{A}}} P^*(x | x_{<i}) \\ + \beta_{\text{sup}} \sum_{x \in S_{x_{<i}}^* \cap \mathcal{A}} P_{\text{trunc}}(x | x_{<i}) \end{aligned} \quad (4)$$

The first term represents the total probability mass of the true distribution lost to truncation, weighted by hyperparameter β_{var} . The second term represents the total probability mass placed off the support of the true distribution (thus constituting a bad continuation), weighted by β_{sup} .⁵

Since the mass of a word under the true model, $P^*(x | x_{<i})$, may be arbitrarily close to zero, it is hard to guarantee that the first term (β_{var}) is zero. One cannot guarantee that any non-complete allowed set \mathcal{A} contains the full support of P^* . However, the smoothing model does provide bounds on the probabilities of words in $S_{x_{<i}}^* \cap \mathcal{A}$, meaning we can in principle avoid unnecessarily truncating words while still maintaining zero cost from the β_{sup} precision term. While we cannot know the exact properties of the unobserved smoothing distribution, we can use this fact to design principles desmoothing algorithms should follow.

3.4 Principles for truncation as desmoothing

Our LM framing specifies bounds on the probabilities of words outside the support of the true distribution, and our TV_S motivates minimizing the difference between the allowed set $\mathcal{A}_{x_{<i}}$ and the support $S_{x_{<i}}^*$. We now use both of these to describe principles for truncation; if these principles are not met, the word is in the support of $S_{x_{<i}}^*$ and should not be truncated.

Absolute probability. Under our smoothing model (Section 3.2), a word outside the support of $P^*(X_i | x_{<i})$ has a bound on its probability:

$$\max_{x \notin S_{x_{<i}}^*} P_\theta(x | x_{<i}) \leq (1 + \delta)(1 - \bar{\lambda})/|\mathcal{V}|, \quad (5)$$

since we posited that smoothing never accounts for more than $\bar{\lambda}$ of the distribution. While these terms are not known, the bound is likely small (since δ is small). Hence as a general principle, words with

large probability should not be truncated, since above a small probability threshold, they must be in the support of P^* .

Relative probability. Under our model, a distribution with high entropy has less smoothing, that is, λ is smaller, e.g., note the term $\exp(-h_{x_{<i}})$ in the bound on λ . This directly results in a lower maximum probability a word outside the support of the true distribution can achieve:

$$\max_{x \notin S_{x_{<i}}^*} P_\theta(x) \leq \alpha \exp(-h), \quad (6)$$

where $\exp(-h_{x_{<i}})$ is the probability of a word in the uniform distribution of entropy $h_{x_{<i}}$ (and α is a constant). The general principle is to only truncate words whose probabilities are also low relative to the rest of the distribution.

In uniform dist.
 $h(x) = -\sum \frac{1}{N} \log \frac{1}{N}$
 $= \log N$
 $\Rightarrow N = \exp(h)$
 $\frac{1}{N} = \exp(-h)$

3.5 Desmoothing and n -gram models

The issue of smoothing on sample quality is apparent in n -gram language models. An n -gram language model MLE estimate explicitly counts the number of times each $(n-1)$ -word phrase is followed by a word in \mathcal{V} . To avoid infinite perplexity (as the count estimates are zero almost everywhere), an n -gram model is explicitly smoothed (Katz, 1987; Church and Gale, 1991).

Text generated from *unsmoothed* n -gram models is locally coherent.⁶ However, we show that n -gram models smoothed with the uniform distribution generate nonsense (Figure 2). Why is this? Consider a 5-gram LM smoothed with the uniform distribution. If x' is sampled from outside the support of the 5-gram model's support, then the new history (x_{i-1}, x') was never seen during the training of the 5-gram model, so now the model has only the poorly estimated probabilities from the smoothing distribution.

4 Methods

We now describe in detail two popular truncation sampling algorithms, discuss how they break our desmoothing principles, and then present two new truncation sampling algorithms including our proposed η -sampling.

4.1 Top- p (nucleus) sampling

Top- p (nucleus) sampling truncates words that are outside the minimal set of (most probable)

⁶As noted by Yoav Goldberg <https://nbviewer.org/gist/yoavg/d76121dfde2618422139> and Jurafsky and Martin (2000), Chapter 3: N -gram Language Models.

⁵See Section A.1 for the relationship to the total variation distance.

For words outside the support of true model.

words that account for at least p percent of the distribution. That is, the allowed set is as follows. Let $x^{(1)}, \dots, x^{(|\mathcal{V}|)}$ be the words in \mathcal{V} sorted in order of decreasing probability under $P_\theta(X | x_{<i})$. Then let j be the integer such that $j = \arg \min_{j'} \sum_{i=1}^{j'} P_\theta(x^{(i)} | x_{<i}) \geq p$. The allowed set of top- p sampling is then $\mathcal{A}_{x_{<i}} = \{x^{(1)}, \dots, x^{(j)}\}$.⁷ Top- p sampling breaks the absolute probability principle: words with up to $(1 - p)$ probability may be truncated simply because other high-probability words cover probability p . For the prompt *My name*, the word *is* is assigned 0.96 probability by GPT-2, but less likely candidates *'s*, *was* and *isn* shouldn't be truncated. Intuitively, $(1 - p)$, e.g., 0.05 or 0.01 is quite high probability given a vocabulary size of, e.g., 50,000.

4.2 Typical decoding

Typical decoding is motivated by local informativeness: never generate words that are too surprising or too predictable (Meister et al., 2022a). The algorithm sorts the vocabulary in order of the difference between the entropy $h_{\theta, x_{<i}}$ of the LM conditional distribution and the negative log-probability of the word, and takes words from this list to cover p percent of the distribution. That is, let $x^{(1)}, \dots, x^{(|\mathcal{V}|)}$ be the words in \mathcal{V} in sorted order of increasing $|h_{\theta, x_{<i}} + \log p_\theta(x | x_{<i})|$.⁸ Then let j be the integer $j = \arg \min_{j'} \sum_{i=1}^{j'} P_\theta(x^{(i)} | x_{<i}) \geq p$. The allowed set of typical decoding is $\mathcal{A}_{x_{<i}} = \{x^{(1)}, \dots, x^{(j)}\}$. This breaks the absolute probability principle for the same reason as top- p , and additionally can truncate the most probable words.

4.3 ϵ -sampling (ours)

The absolute probability principle—that words outside the support of the true distribution have low probability—suggests a simple truncation algorithm: for some hyperparameter threshold ϵ allow any word with greater than ϵ probability.

$$\mathcal{A}_{x_{<i}} = \{x \in \mathcal{V} : P_\theta(x | x_{<i}) > \epsilon\} \quad (7)$$

In the case of the prompt *My name* where top- p rejects plausible words because of the probability assigned to *is* (and *'s*), ϵ -sampling allows additional words with a threshold of, e.g., 0.0003.

However, ϵ -sampling breaks the relative probability principle. For example, the prompt *The* should allow many continuations, and top- p with

⁷Often, p is taken as 0.9 or 0.95.

⁸ $h_{\theta, x_{<i}} = -\sum_{x \in \mathcal{V}} P_\theta(x | x_{<i}) \log P_\theta(x | x_{<i})$.

GPT-2 allows over ten thousand words, but ϵ would have to be impractically small to do so. This is a key failure akin to that of top- k sampling; when many next words are plausible, the allowed set should reflect that.

4.4 η -sampling (ours)

Our proposed algorithm, η -sampling, composes respect for both the absolute and relative probability principles. Consider a conditional distribution $P_\theta(X | x_{<i})$ with entropy $h_{\theta, x_{<i}}$. The probability of a word in the uniform distribution of entropy $h_{\theta, x_{<i}}$ is $\exp(-h_{\theta, x_{<i}})$. Our entropy-dependent threshold is $\alpha \exp(-h_{\theta, x_{<i}})$ where $\alpha \in [0, 1]$. Combining this rule with our epsilon rule for the absolute probability principle, we come to:

$$\mathcal{A}_{x_{<i}} = \{x \in \mathcal{V} | P_\theta(x | x_{<i}) > \eta\}$$

$$\eta = \min(\epsilon, \alpha \exp(-h_{\theta, x_{<i}}))$$

where $h_{\theta, x_{<i}}$ is the entropy of $P_\theta(X | x_{<i})$. In this work, to expose a single hyperparameter, we set $\alpha = \sqrt{\epsilon}$, which we find works well empirically. For intuition, think of $\epsilon \approx 0.0009$.

Analysis of η -sampling. Returning to our smoothing model, we note that η -sampling approximates optimal desmoothing in the regime that the support penalty β_{sup} dominates the variation penalty β_{var} . Consider a truncation algorithm that truncates as η -sampling, but sets η as:

$$\eta = \min\left(\frac{(1 - \bar{\lambda})(1 + \delta)}{V}, \alpha \exp(-h_{x_{<i}})\right), \quad (8)$$

where $h_{x_{<i}}$ is the entropy of the true distribution, not P_θ . We're guaranteed that the support loss (the term weighted by β_{sup}) is zero, and that the variation loss (weighted by β_{var}) is minimized relative to the constraint of zero support loss. If $x \notin S_{x_{<i}}^*$, then the probability of x is less than or equal to the min of $(1 - \bar{\lambda})(1 + \delta)/V$ and $\frac{V \alpha \exp(-h_{x_{<i}})}{1 + \delta} \times \frac{1 + \delta}{V} = \alpha \exp(-h_{x_{<i}})$. So, we're guaranteed that $\mathcal{A}_{x_{<i}} \subseteq S_{x_{<i}}^*$, and truncating more would break this guarantee.⁹ Our η -sampling approximates this by using the LM entropy instead of the unavailable true distribution entropy, and without knowing the true hyperparameters.

5 Experiments & Results

Our experiments characterize η -sampling relative to the state-of-the-art top- p and typical decoding.

⁹See Appendix A.2 for an expanded version of this argument.

Since 0.05 or 0.1 is a high prob. given sok vocab size, we are getting rid of high prob values in top-p and typical sampling.

Method	Hyperparameters
top- p	{0.89, 0.9, 0.92, 0.95, 0.99}
typical	{0.2, 0.9, 0.92, 0.95}
ϵ	{0.001, 0.0009, 0.0006, 0.0003, 0.0001}
η	{0.004, 0.002, 0.0009, 0.0006, 0.0003}

Table 1: Hyperparameter sweep for each method.

Method \ Model	sm	med	lg	xl
raw sampling \dagger	0.589	0.373	0.845	0.882
top- p \dagger	0.878	0.915	0.936	0.940
top- p (our replication)	0.874	0.917	0.932	0.944
Typical Decoding	0.873	0.906	0.922	0.939
ϵ -sampling (ours)	0.874	0.918	0.936	0.941
η -sampling (ours)	0.880	0.920	0.935	0.942

Table 2: Results on the MAUVE metric for open-ended GPT-2 WebText generation. Higher is better. The \dagger indicates numbers drawn from Pillutla et al. (2021). Bold indicates best for model, not necessarily significantly.

We use MAUVE, an automatic metric for open-ended generation, to find hyperparameters giving comparable diversity-accuracy tradeoffs. η -sampling behaves better in a range of settings, from long-document generation to more defensibly truncating low-entropy distributions.

Models & Data. In all experiments, we use all or some subset of the four GPT-2 models (Radford et al., 2019) of varying sizes. Experiments are run on in-distribution, held-out data from the validation or test set of GPT-2 (WebText), since it is composed of a wide variety of long-form documents.

5.1 Hyperparameter sweep on MAUVE

We first find hyperparameters for each of top- p , typical decoding, ϵ -sampling, and η -sampling that maximize MAUVE score for each GPT-2 model on WebText.

Setting. Following the MAUVE paper’s setting exactly (Pillutla et al., 2021), we take the GPT-2 family of models and 5,000 samples from their test data. For each sample, we prompt the model with 35 words and generate until at most 1024 words. We study GPT-2 small (124M parameters), medium (355M), large (774M) and XL (1.5B) models.

Evaluation. MAUVE attempts to measure both the precision (are samples generally like those from the true distribution) and recall (is the variability in samples like that of those from the true distribution)

Study 1: Human vs top- p vs η			
	top- p	η -sampling	Human
Top- p vs human	43 (43%)	—	56 (56%)
η vs human	—	42 (42%)	53 (53%)
Top- p vs η	39 (39%)	53 (53%)	—

Study 2: top- p vs η -sampling			
	Top- p	η -sampling	Equal
Top- p vs η	118 (40%)	159 (53%)	17 (6%)

Table 3: Human preferences of long-document plausibility; we report absolute numbers of judgments, and percentages in parentheses. Judgment percents that both suffixes were too bad to judge can be inferred.

of samples from a text generation system. It was shown by Pillutla et al. (2021) to correlate well with human judgments.

Hyperparameters. Top- p , typical decoding, ϵ -sampling, and η -sampling all have a hyperparameter which determines the severity of truncation. The set we search over is given in Table 1.¹⁰ We pick the best hyperparameter using 2–5 seeds on the validation set, and report the average performance across 5 seeds on the test set.

Results. The results are reported in Table 2; we find that overall, the methods perform similarly, with typical decoding performing slightly worse than top- p and our methods.

5.2 Human evaluation of long-document suffix plausibility

We now study whether η -sampling leads to more coherent long-document generations than top- p sampling. We omit typical decoding since it does not seem to outperform top- p on MAUVE. Considering that holistic evaluation of long texts is difficult for humans (Ippolito et al., 2020) we design a human study to evaluate *long document plausibility*: given a shared document prefix, which method’s generated suffix (omitting the middle) is more reasonably from the same document? This new evaluation avoids forcing humans to keep up to 1024 words in working memory.

Setting. For each of top- p and η -sampling, we sample from GPT-2 large with MAUVE-maximizing hyperparameters, conditioned on each prefix of 35 subword tokens from the WebText validation set. From this set we filter to prefixes for

¹⁰The hyperparameter set for our methods was chosen to have similar average total variation values between pre- and post-truncation to the top- p set.

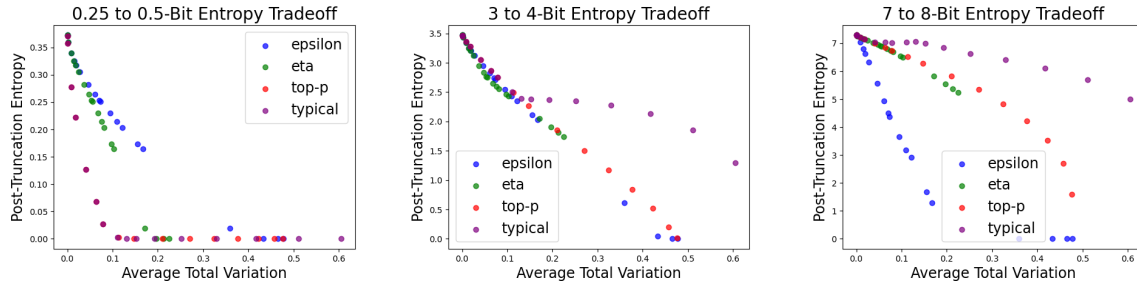


Figure 3: Top- p sampling aggressively truncates low-entropy distributions and ϵ -sampling aggressively truncates high-entropy distributions, while η -sampling strikes a balance.

which the reference and both generated documents are at least 900 tokens long and pass manual filter for quality.¹¹ 59 workers from the United States were recruited on Amazon Mechanical Turk with the Master qualification, and paid \$1 per task with an expected time of 3.5 to 4 minutes. We run two studies.

Study 1. We show a human evaluator the 35-token prefix, as well as the last 70 tokens of two documents (of the 3 possible). The evaluator is asked to judge which of the two suffixes may *more reasonably be from the same document* as the prefix, or to note that both are too bad to judge. For each of the three possible pairings of top- p , η -sampling, and reference document, we elicit 100 human judgments over 100 prefixes.

Study 2. We ran a second study just comparing top- p to η -sampling to allow for larger n , since we had finite resources and the result that both methods generate text worse than humans is not at issue. To test whether the effect size observed was in part due to forcing evaluators to pick one of the two methods, in this study we allow human evaluators to mark that both suffixes are of equal quality.

Results. The results are reported in Table 3. In Study 1, we find that human document generations are preferred over top- p and η -sampling at roughly the same rate, while η -sampling is preferred over top- p (53% to 40%). In Study 2, we find that η -sampling is significantly preferred more frequently than top- p with a Wilcoxon paired test ($p = 0.0138$) at the same effect size.

5.3 Entropy analysis

We now want to build a deeper understanding of the characteristics of the algorithms: what parts of the

distribution tend to get cut by each method? In our first analysis, we study whether each method has a tendency to aggressively truncate distributions of a given entropy. A low-entropy distribution might be given by the prompt *Barack Obama went to the White ...*, while a high-entropy distribution might be given by the prompt *My name is ...*

Setting. For a range of hyperparameters, we plot the average amount of truncation across all contexts against the retained entropy for an entropy range. We use total variation to measure average truncation, $\mathbb{E}_{x_{<i} \sim P(X)} \|P_\theta(X_i | x_{<i}) - P_{\text{trunc}}(X_i | x_{<i})\|_{\text{TV}}$. For each entropy range R , we consider the set \mathcal{X}_R of prefixes $x_{<i}$ with pre-truncation entropy $h_{\theta, x_{<i}}$ in R and compute the average remaining entropy $\frac{1}{|\mathcal{X}_R|} \sum_{\mathcal{X}_R} h_{\text{trunc}, x_{<i}}$ after truncation.

Results. The results for GPT-2 XL are presented in Figure 3. We find that top- p sampling heavily truncates low-entropy distributions compared to ϵ -sampling and η -sampling. ϵ -sampling heavily truncates high-entropy distributions. Typical behaves like top- p for low-entropy distributions, and retains more entropy in high-entropy distributions.¹² η -sampling strikes a good balance of not heavily truncating low- or high-entropy distributions.

5.4 Repetition analysis

We hypothesize that the tendency of top- p sampling to heavily truncate low-entropy distributions causes it to generate repetitive text by only allowing the repetition-continuing word. To stress test the methods, we devise an adversarial setting in which the prompt has repetitions (as may be the case due to noisy input or natural repetition) and then determine whether the methods break the repetition.

¹¹We also manually filter prompts for quality, following Pillutla et al. (2021). See Appendix B.3.

¹²This is likely because typical decoding cuts the non-uniform head of the distribution, and keeps the more-uniform middle.

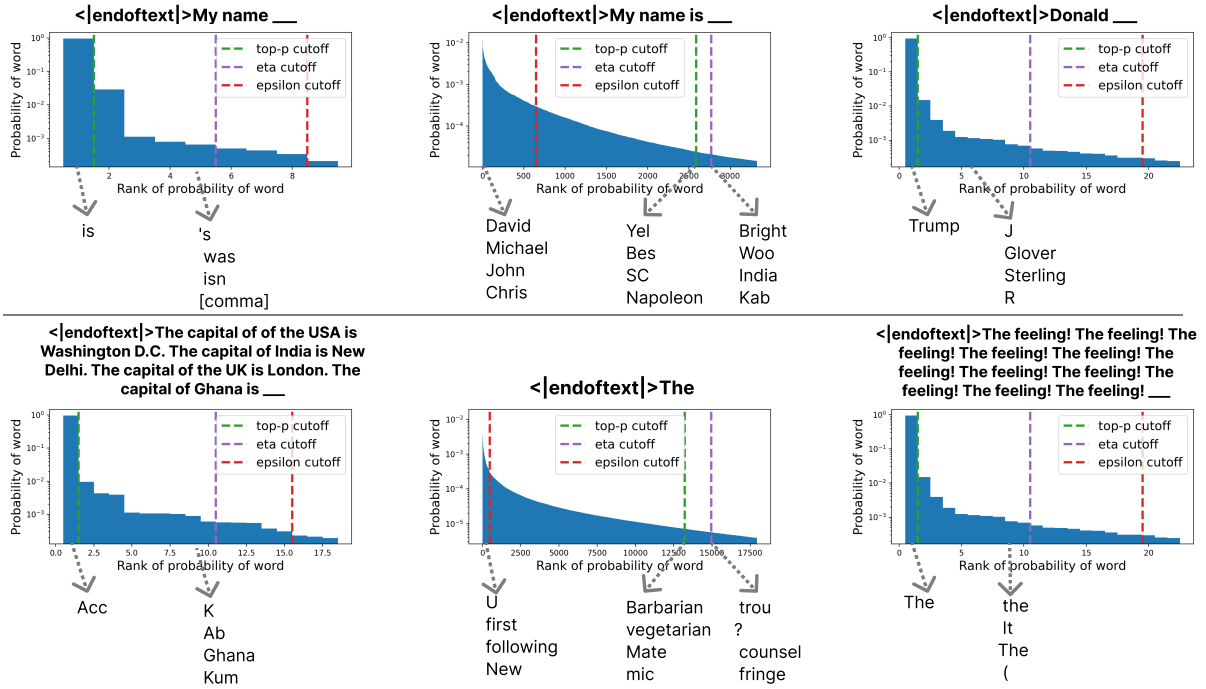


Figure 4: Unit tests of the truncation behavior of top- p , ϵ , and η -sampling on CheckList-inspired prefixes.

Truncation \ Model	Repetition Percent			
	sm	med	lg	xl
top- p	54%	61%	47%	27%
typical	51%	61%	56%	37%
ϵ -sampling (ours)	28%	37%	23%	11%
η -sampling (ours)	37%	40%	26%	12%

Table 4: Table showing repetition-degeneration rates for each method in an adversarial setting; lower is better.

Setting. We take natural prompts—the first 35 words of the Wikipedia biographies of the 101 people with the most-read Wikipedia pages—and synthetically corrupt them by repeating the last 3 subword tokens 5 additional times. Even with the existing repetition in the prompt, we want models to break the cycle and generate normal text again. Here’s an example prompt:

Shawn Corey Carter (born December 4, 1969), known professionally as Jay-Z, is an American rapper, songwriter, record executive, entrepreneur, and media proprietor and media proprietor and media proprietor and media proprietor and media proprietor and media proprietor and media proprietor

For each prompt, we generate 5 completions of up to 512 words. For each of the GPT-2 models, we take the hyperparameter for each truncation sampling algorithm from Section 5.1, and compute the

percent of completions that continue to repeat.¹³

Results. ϵ -sampling achieves the lowest repetition rate, with e.g., 23% for GPT-2 large, while η -sampling performs slightly worse (e.g., 26%). Top- p causes considerably more repetition (e.g., 47%). Typical sampling causes slightly more repetition than top- p .¹⁴

5.5 Studying individual distributions

We now study specific truncation decisions made by each algorithm, to provide more detailed behavioral insights. We construct prompts and observe the truncation behavior of each algorithm on the resulting distribution, treating each as a CheckList-like unit test (Ribeiro et al., 2020).

Setting. We take the GPT-2 large model, provide it with each of 6 prompts, and using the MAUVE-maximizing hyperparameters we found in Section 5.1, truncate the resulting distribution. The prompts are shown in Figure 4. For this experiment we only study top- p , ϵ , and η -sampling.

¹³Any sample with less than 1 average negative log probability under the model is labeled a repetition. We found this more useful than n -gram repetition statistics, as, e.g., repetition can involve small variation.

¹⁴This is likely because the MAUVE-maximizing hyperparameter for typical sampling (e.g., 0.92 for GPT-2 large) is generally more conservative than that for top- p (e.g., 0.95.)

Results. The results are visualized in Figure 4. We use two low-entropy prompts, *My name...* and *Donald...* and in both cases, find that top- p decoding only allows a single word continuation. Top- p can only generate *is* after *My name*, and *Trump* after *Donald*, which we find undesirable; we would like our truncation to allow, e.g., multiple Donalds to be discussed. For a prompt with the phrase *The feeling!* repeated multiple times (as one might say euphorically), top- p can only continue the repetitive pattern, unlike ϵ and η -sampling. For a prompt suggesting specification of capitals of countries, we find that top- p only allows the *correct* capital name, whereas η -sampling and ϵ -sampling allow different continuations which do not follow the in-context trend, suggesting that top- p may be better for generating, e.g., answers to questions. We use two high-entropy prompts, *The...* and *My name is...*, finding that η -sampling and top- p sampling allow a range of possibilities, unlike ϵ -sampling. The behavior of ϵ -sampling in allowing fewer words in higher entropy conditional distributions is a clear failure.

6 Related Work

Stochastic decoding algorithms. Stochastic decoding algorithms produce sequences from a model and involve randomness. The simplest is *sampling*, sometimes called *ancestral sampling*, (Bishop, 2006), which generates a sample from the model. Some stochastic decoding methods attempt to find high-likelihood sequences instead of attempting to recreate the true distribution, like stochastic beam search (Kool et al., 2019) and conditional poisson stochastic beam search (Meister et al., 2021a). Truncation sampling algorithms, like top- k (Fan et al., 2018), top- p (Holtzman et al., 2020), and Mirostat (Basu et al., 2021), are intended to improve quality but keep variety. Welleck et al. (2020) found that truncation algorithms can lead to non-zero mass assigned to infinite sequences.

KL-divergence, language models, smoothing. The most famous example of methods that do not cover every mode is GANs (Goodfellow et al., 2014). In language modeling, some have pointed to the inability of the softmax function to assign 0 probability to any category as a deficiency and proposed sparse alternatives (Martins and Astudillo, 2016; Peters et al., 2019; Tezkebayev et al., 2021). This intuition is akin to ours, as is loss truncation (Kang and Hashimoto, 2020), which keeps rare

events from incurring arbitrarily high loss. Mohri and Roark (2006) attempt to identify structural zeros in the distribution of language when inducing probabilistic context-free grammars.

High-entropy language generation & evaluation. Evaluation of open-ended generation of natural language is difficult; one must evaluate both the quality of samples and the diversity. Quality is hard to measure in high-entropy generation, and is often not correlated with model probability (Hashimoto et al., 2019; Meister et al., 2022b). An emergent line of work connects human notions of quality, and human generative tendencies, with the uniform information density hypothesis (e.g., leading to typical decoding) (Wei et al., 2021; Meister et al., 2021b). Both Meister and Cotterell (2021) and Pillutla et al. (2021) directly estimate whether model samples’ statistics match those of natural language. Nadeem et al. (2020) study properties held by successful strategies for reallocating mass away from the tail of LM distributions.

7 Conclusion

We’ve framed the class of truncation sampling algorithms as performing desmoothing, an insight that led to principles for how truncation should be done to recover the training distribution, a new truncation sampling algorithm, and evaluations that show the deficiencies of existing algorithms. We find the tendency of top- p decoding to over-truncate low-entropy distributions to be particularly surprising. We aim for these insights, and the evaluations we use, to drive further research in understanding and improving how we generate from neural language models.

Acknowledgements

The authors would like to thank John Thickstun, Rishi Bommasani, Kaitlyn Zhou, Will Merrill, Nelson Liu, and Tatsunori Hashimoto for helpful discussions on this work, and to the reviewers for clarifying feedback. JH was supported by an NSF Graduate Research Fellowship under grant number DGE-1656518. We gratefully acknowledge the support of a PECASE Award.

8 Limitations

With the analysis we’ve done, we believe it to be very difficult to derive an understanding of all

the *sequence-level* effects truncation sampling algorithms (including ours) have: what kinds of sequences are we disallowing? What types, or sources of language are being (unknowingly) disallowed? Beyond this, we’ve only tested our algorithms on English language models; the conditional distributions of languages with rich morphology likely have different properties (especially with subword models).

9 Ethics Statement

Any work to improve generative models of text comes with ethical concerns surrounding negative use cases of text generation including hate speech and misinformation. While our algorithm does improve long text generation, we hope it also provides insight into the unintended and until-now unknown consequences of existing truncation sampling algorithms (including top- p). Algorithms like ours, which reallocate probability mass from the least likely elements of a distribution, have a particular risk of harm in removing the ability of models to talk about topics or names that are already rare. Concurrent work finds that the choice of stochastic decoding algorithm affects measured fairness metrics in open-ended generation (Dhamala et al., 2022). Our framing, and the hope for future work, is to use truncation to recover something as close to the *training* distribution as possible; of course, the training distribution must then be chosen with care. Generating a word due to smoothing (noise) would likely mean that subsequently generated words about that topic would be low-quality, which is also undesirable.

References

- Sourya Basu, Govardana Sachitanandam Ramachandran, Nitish Shirish Keskar, and Lav R. Varshney. 2021. [MIROSTAT: A neural text decoding algorithm that directly controls perplexity](#). In *International Conference on Learning Representations*.
- Christopher M Bishop. 2006. *Pattern recognition and machine learning*. Springer.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. 2022. PaLM: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311*.
- Kenneth W Church and William A Gale. 1991. A comparison of the enhanced Good-Turing and deleted estimation methods for estimating probabilities of English bigrams. *Computer Speech & Language*, 5(1):19–54.
- Jwala Dhamala, Varun Kumar, Rahul Gupta, Kai-Wei Chang, and Aram Galstyan. 2022. An analysis of the effects of decoding algorithms on fairness in open-ended language generation. In *2018 IEEE Spoken Language Technology Workshop (SLT)*. IEEE.
- Angela Fan, Yacine Jernite, Ethan Perez, David Grangier, Jason Weston, and Michael Auli. 2019. [ELI5: long form question answering](#). In *Proceedings of the 57th Conference of the Association for Computational Linguistics, ACL 2019, Florence, Italy, July 28- August 2, 2019, Volume 1: Long Papers*, pages 3558–3567. Association for Computational Linguistics.
- Angela Fan, Mike Lewis, and Yann Dauphin. 2018. [Hierarchical neural story generation](#). In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 889–898, Melbourne, Australia. Association for Computational Linguistics.
- Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, Shawn Presser, and Connor Leahy. 2021. [The Pile: An 800gb dataset of diverse text for language modeling](#). *CoRR*, abs/2101.00027.
- Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. 2014. Generative adversarial nets. *Advances in neural information processing systems*, 27.
- Tatsunori B Hashimoto, Hugh Zhang, and Percy Liang. 2019. Unifying human and statistical evaluation for natural language generation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 1689–1701.
- Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. 2020. [The curious case of neural text de-generation](#). In *International Conference on Learning Representations*.
- Daphne Ippolito, Daniel Duckworth, Chris Callison-Burch, and Douglas Eck. 2020. [Automatic detection of generated text is easiest when humans are fooled](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1808–1822, Online. Association for Computational Linguistics.

- Daniel Jurafsky and James H. Martin. 2000. *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*, 1st edition. Prentice Hall PTR, USA.
- Daniel Kang and Tatsunori B. Hashimoto. 2020. [Improved natural language generation via loss truncation](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 718–731, Online. Association for Computational Linguistics.
- S. Katz. 1987. [Estimation of probabilities from sparse data for the language model component of a speech recognizer](#). *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 35(3):400–401.
- Wouter Kool, Herke Van Hoof, and Max Welling. 2019. [Stochastic beams and where to find them: The Gumbel-top-k trick for sampling sequences without replacement](#). In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 3499–3508. PMLR.
- Yang Liu and Mirella Lapata. 2019. [Text summarization with pretrained encoders](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3730–3740, Hong Kong, China. Association for Computational Linguistics.
- Andre Martins and Ramon Astudillo. 2016. From softmax to sparsemax: A sparse model of attention and multi-label classification. In *International conference on machine learning*, pages 1614–1623. PMLR.
- Clara Meister, Afra Amini, Tim Vieira, and Ryan Cotterell. 2021a. [Conditional Poisson stochastic beams](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 664–681, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Clara Meister and Ryan Cotterell. 2021. [Language model evaluation beyond perplexity](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 5328–5339, Online. Association for Computational Linguistics.
- Clara Meister, Tiago Pimentel, Patrick Haller, Lena Jäger, Ryan Cotterell, and Roger Levy. 2021b. [Revisiting the Uniform Information Density hypothesis](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 963–980, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Clara Meister, Tiago Pimentel, Gian Wiher, and Ryan Cotterell. 2022a. [Typical decoding for natural language generation](#). *CoRR*, abs/2202.00666.
- Clara Meister, Gian Wiher, Tiago Pimentel, and Ryan Cotterell. 2022b. [On the probability–quality paradox in language generation](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 36–45, Dublin, Ireland. Association for Computational Linguistics.
- Mehryar Mohri and Brian Roark. 2006. [Probabilistic context-free grammar induction based on structural zeros](#). In *Proceedings of the Human Language Technology Conference of the NAACL, Main Conference*, pages 312–319, New York City, USA. Association for Computational Linguistics.
- Moin Nadeem, Tianxing He, Kyunghyun Cho, and James Glass. 2020. A systematic characterization of sampling algorithms for open-ended language generation. In *Proceedings of the 1st Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 10th International Joint Conference on Natural Language Processing*, pages 334–346.
- Ian Osband, Zheng Wen, Seyed Mohammad Asghari, Vikranth Dwaracherla, Morteza Ibrahimi, Xiuyuan Lu, and Benjamin Van Roy. 2022. Epistemic neural networks. *arXiv preprint arXiv:2107.08924*.
- Ben Peters, Vlad Niculae, and André FT Martins. 2019. Sparse sequence-to-sequence models. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 1504–1519.
- Krishna Pillutla, Swabha Swayamdipta, Rowan Zellers, John Thickstun, Sean Welleck, Yejin Choi, and Zaid Harchaoui. 2021. [Mauve: Measuring the gap between neural text and human text using divergence frontiers](#). In *Advances in Neural Information Processing Systems*, volume 34, pages 4816–4828. Curran Associates, Inc.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. 2019. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9.
- Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. [Beyond accuracy: Behavioral testing of NLP models with CheckList](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4902–4912, Online. Association for Computational Linguistics.
- Maxat Tezkbayev, Vassilina Nikoulina, Matthias Gallé, and Zhenisbek Assylbekov. 2021. [Speeding up entmax](#). *CoRR*, abs/2111.06832.

Jason Wei, Clara Meister, and Ryan Cotterell. 2021. [A cognitive regularizer for language modeling](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 5191–5202, Online. Association for Computational Linguistics.

Sean Welleck, Ilia Kulikov, Jaedeok Kim, Richard Yuanzhe Pang, and Kyunghyun Cho. 2020. Consistency of a recurrent language model with respect to incomplete decoding. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 5553–5568.

A Notes

A.1 Support-weighted total variation

We introduce new notation just for this section, to present support-weighted total variation in generality. Recall that the total variation distance between discrete distribution R over space \mathcal{V} and discrete distribution U_t , the result of truncation with allowed set $\mathcal{A} \subseteq \mathcal{V}$ from a discrete distribution U over \mathcal{V} , is

$$\sum_{x \in \mathcal{V}} |R(x) - U_t(x)|. \quad (9)$$

Denoting the support of R as S_R , we can partition \mathcal{V} into four sets:

$$\begin{aligned} S_R \cap \bar{\mathcal{A}} \\ \bar{S}_R \cap \mathcal{A} \\ S_R \cap \mathcal{A} \\ \bar{S}_R \cap \bar{\mathcal{A}} \end{aligned} \quad (10)$$

We split the sum of the total variation distance into these four terms.

The first represents the words that are in the support of R but not in the allowed set of U_t :

$$\sum_{S_R \cap \bar{\mathcal{A}}} |R(x) - U_t(x)| = \sum_{S_R \cap \bar{\mathcal{A}}} R(x), \quad (11)$$

since $U_t(x) = 0$ if $x \notin \mathcal{A}$. This exactly represents the total probability mass that was lost from R . The second term represents the words that are not in the support of R but were allowed:

$$\sum_{\bar{S}_R \cap \mathcal{A}} |R(x) - U_t(x)| = \sum_{\bar{S}_R \cap \mathcal{A}} U_t(x), \quad (12)$$

since $R(x) = 0$ if $x \notin S_R$. This exactly represents the total probability that we sample a word from U_t

that has zero probability under R (and so we move off the support of R for future generation.) the third term is the words that were correctly allowed:

$$\sum_{S_R \cap \mathcal{A}} |R(x) - U_t(x)|. \quad (13)$$

In this case, $U_t(x)$ may be an under or overestimate of $R(x)$. The last term is the words that were correctly truncated:

$$\sum_{\bar{S}_R \cap \bar{\mathcal{A}}} |R(x) - U_t(x)| = \sum_{\bar{S}_R \cap \bar{\mathcal{A}}} |0 - 0| \quad (14)$$

which is identically zero.

To form our support-weighted total variation metric, we took the first two terms, which are interpretable and each exactly specifies one of the two desiderata from a truncation algorithm: maintaining the variety of R , and not generating a word that R wouldn't generate. However, in different use cases, one or the other may be more crucial; hence we give each its own hyperparameter, β_{var} and β_{sup} , to arrive at our metric,

$$\begin{aligned} \text{TV}_S(R, U_t) = & \beta_{\text{var}} \sum_{x \in S_R \cap \bar{\mathcal{A}}} R(x) \\ & + \beta_{\text{sup}} \sum_{x \in \bar{S}_R \cap \mathcal{A}} U_t(x). \end{aligned} \quad (15)$$

A.2 Analysis of η -sampling

The purpose of this analysis is to show that if one assumes our smoothing model, then an η -sampling approximates an algorithm that avoids sampling from outside the support of the true distribution while minimally truncating the distribution.

Consider a conditional distribution from a language model under our model, $P_\theta(X_i | x_{<i})$. Consider an allowed set $\mathcal{A}_{x_{<i}}$ defined via a probability threshold, $\mathcal{A} = \{x | P_\theta(x | x_{<i}) > \eta^*\}$, where η^* is defined as

$$\eta^* = \min \left(\frac{(1 - \bar{\lambda})(1 + \delta)}{V}, \alpha \exp(-h_{x_{<i}}) \right). \quad (16)$$

In this case, it is guaranteed that $x \in S_{x_{<i}}^*$, since η^* represents the maximum probability of a word whose probability stems entirely from the smoothing distribution.

If one sets a lower probability threshold $\eta' = \eta^* - \psi$ for some $\psi > 0$ when computing the allowed set, then under our model, there can be a conditional distribution such that $x \notin S_{x_{<i}}^*$, and $P_\theta(x | x_{<i}) > \eta'$. Such an x would be incorrectly allowed.

Method \ Model	small	med	large	XL
Top- p	0.9	0.89	0.95	0.95
Typical	0.9	0.9	0.92	0.92
ϵ -sampling	0.0006	0.0009	0.0003	0.0003
η -sampling	0.002	0.0006	0.0006	0.0003

Table 5: Best-performing hyperparameters according to MAUVE from experiments in Section 5.1.

Similarly, if one sets a higher probability threshold $\eta' = \eta^* + \psi$ for some $\psi > 0$ when computing the allowed set, then under the model, there can be a conditional distribution such that $x \in S_{x_{<i}}^*$, and $P_\theta(x \mid x_{<i}) \in (\eta, \eta')$. Defining the allowed set with η' , we truncate x , which is unnecessary, since words in $S_{x_{<i}}^*$ have probability at least η under the language model.

This argument has considered truncation algorithms that specify their allowed set as every word in \mathcal{V} with LM probability above a threshold, showing that setting the threshold as η^* guarantees (under our model) that we sample from the support of the true distribution without unnecessarily truncating too much. We now consider allowed set defined by algorithms other than probability thresholds. Let the allowed set defined according to the η^* threshold be $\mathcal{A}_{x_{<i}}^*$. Consider an allowed set $\mathcal{A}_{x_{<i}}$ defined by another truncation sampling algorithm (which may not define it via a probability threshold like. If $\mathcal{A}_{x_{<i}} = \mathcal{A}_{x_{<i}}^*$, then the two algorithms are indistinguishable for this prefix. Otherwise, if $x \in \mathcal{A}_{x_{<i}}$ and $x \notin \mathcal{A}_{x_{<i}}^*$, then x may be outside the support of the true distribution, and should have been truncated. And if $x \in \mathcal{A}_{x_{<i}}^*$ and $x \notin \mathcal{A}_{x_{<i}}$, then x was unnecessarily truncated.

When using our η -sampling algorithm, we neither know the true hyperparameters, nor do we have access to the true distribution conditional entropy, so η -sampling only approximates this. Specifically, we set the hyperparameters of η -sampling via search on the task of interest, and we use the observed LM entropy instead of the true distribution entropy in computing the relative probability threshold. In practice, one wants to set a threshold of truncation based on the needs of the task and the tolerance for error, so a threshold that perfectly excludes words outside the true distribution support may not be optimal for the task of interest anyway.

B More Experimental Details

B.1 Hyperparameters

The MAUVE-maximizing hyperparameters for each truncation sampling algorithm for each model are provided in Table 5.

B.2 5-gram model

For our small demo demonstrating the behavior of smoothed n -gram models, we trained a 5-gram model on 10,000 documents from The Pile (Gao et al., 2021). We smoothed the model with the uniform distribution.

B.3 Amazon Mechanical Turk Details

To provide more transparency into our human studies, we provide the form that was shown to human annotators for both of our studies. The (similar) interfaces shown for Study 1 and Study 2 are shown in Figure 5 and Figure 6, respectively. We randomize the ordering of presentation of the methods' generations (note that the forms say "Option 1" and "Option 2".)

Of the 59 unique workers, 44 unique workers participated in study 1, and 36 unique workers participated in study 2.

We follow Pillutla et al. (2021) in manually filtering the WebText prompts that go into our human study. Webtext is noisy, and not all prompts are clearly natural language. Our manual filtering of prompts led to 36 rejected prompts (of 146 considered) due to quality for study 1. Our manual filtering of prompts led to 100 rejected prompts (of 402 considered) due to quality for study 2. This is compared to rejecting 3169 of 5000 prompts due to quality in the original MAUVE paper; we attempted to minimally filter while guaranteeing that prompts were natural language. Our kept and filtered prompts are available in our codebase.

- You are given the beginning of a document that appeared on the web.
- Two options are given for possible passages that are intended to be from the same document, but occur long after the beginning (i.e., hundreds of words are omitted between the beginning and the option passage.)
- Your task is to **pick the passage option that seems more reasonably could be in the same document as the beginning paragraph**.
- It is possible that both passage options look comparable, or both are very bad. However, you should try to discern carefully and pick the better one between the two.
- When judging the plausibility of the passage options, please **do not consider** the fact that each passage may start or end awkwardly (i.e., in the middle of a word) but please **do consider** whether it seems to be about the same or similar subjects as the beginning paragraph, mentions the same or similar people or things as the beginning paragraph, has the same or similar style as the beginning paragraph, and makes sense within itself using your best judgement.

Beginning of document: \${prefix}	
Option 1 for passage from same document: \${option_1}	Option 2 for passage from same document: \${option_2}

Option 1	Option 2
<input type="radio"/> Option 1 more plausibly comes from the same document as the beginning paragraph.	<input type="radio"/> Option 2 more plausibly comes from the same document as the beginning paragraph.

☐ I can't make sense of either option.

Submit

Figure 5: The interface shown to human annotators for Study 1.

- you are given the beginning of a document that appeared on the web.
- Two options are given for possible passages that are intended to be from the same document, but occur long after the beginning (i.e., hundreds of words are omitted between the beginning and the option passage.)
- Your task is to **pick the passage option that seems more plausibly could be in the same document as the beginning paragraph**.
- It is possible that both passage options look comparable, or both are very bad. You should discern carefully and determine which is better, but can mark that they're "equally plausible" or that you "Can't make sense of either option".
- When judging the plausibility of the passage options, please **do not consider** the fact that each passage may start or end awkwardly (i.e., in the middle of a word) but please **do consider** whether it seems to be about the same or similar subjects as the beginning paragraph, mentions the same or similar people or things as the beginning paragraph, has the same or similar style as the beginning paragraph, and makes sense within itself using your best judgement.

Beginning of document: \${prefix}		
Option 1 for passage from same document: \${option_1}	Option 2 for passage from same document: \${option_2}	

Option 1	Both equally plausible	Option 2
<input type="radio"/> Option 1 more plausibly comes from the same document as the beginning paragraph.	<input type="radio"/> They're equally plausible. (Avoid marking this if possible)	<input type="radio"/> Option 2 more plausibly comes from the same document as the beginning paragraph.

☐ I can't make sense of either option.

Submit

Figure 6: The interface shown to human annotators for Study 2.